

FLOATING POINT ARITHMETIC - ERROR ANALYSIS

- Brief review of floating point arithmetic
- Model of floating point arithmetic
- Notation, backward and forward errors
- Read: Section 2.7 of text.

Roundoff errors and floating-point arithmetic

➤ The basic problem: The set A of all possible representable numbers on a given machine is finite - but we would like to use this set to perform standard arithmetic operations $(+, *, -, /)$ on an infinite set. The usual algebra rules are no longer satisfied since results of operations are rounded.

➤ Basic algebra breaks down in floating point arithmetic.

Example: In floating point arithmetic.

$$a + (b + c) \neq (a + b) + c$$

✎ Matlab experiment: For 10,000 random numbers find number of instances when the above is true. Same thing for the multiplication..

3-2

Csci 5304 – September 16, 2013

Machine precision - machine epsilon

➤ When a number x is very small, there is a point when $1 + x == 1$ in a machine sense. The computer no longer makes a difference between 1 and $1 + x$.

Definition: the machine epsilon is the smallest number ϵ such that

$$fl(1 + \epsilon) \neq 1$$

This number is denoted by \underline{u} – sometimes by ϵ_{ps} .

✎ Matlab experiment: find the machine epsilon on your computer.

➤ Many discussions on what conditions/ rules should be satisfied by floating point arithmetic. The IEEE standard is a set of standards adopted by many CPU manufacturers.

Rule 1.

$$fl(x) = x(1 + \epsilon), \quad \text{where } |\epsilon| \leq \underline{u}$$

Rule 2.

For all operations \odot (one of $+, -, *, /$)

$$fl(x \odot y) = (x \odot y)(1 + \epsilon_{\odot}), \quad \text{where } |\epsilon_{\odot}| \leq \underline{u}$$

Rule 3.

For $+, *$ operations

$$fl(a \odot b) = fl(b \odot a)$$

✎ Matlab experiment: Verify experimentally Rule 3 with 10,000 randomly generated numbers a_i, b_i .

3-3

Csci 5304 – September 16, 2013

3-4

Csci 5304 – September 16, 2013

Example: Consider the sum of 3 numbers: $y = a+b+c$.

➤ Done as $fl(fl(a+b) + c)$

$$\begin{aligned}\eta &= fl(a+b) = (a+b)(1+\epsilon_1) \\ y_1 &= fl(\eta+c) = (\eta+c)(1+\epsilon_2) \\ &= [(a+b)(1+\epsilon_1) + c](1+\epsilon_2) \\ &= [(a+b+c) + (a+b)\epsilon_1](1+\epsilon_2) \\ &= (a+b+c) \left[1 + \frac{a+b}{a+b+c}\epsilon_1(1+\epsilon_2) + \epsilon_2 \right]\end{aligned}$$

So disregarding the high order term $\epsilon_1\epsilon_2$

$$\begin{aligned}fl(fl(a+b) + c) &= (a+b+c)(1+\epsilon_3) \\ \epsilon_3 &\approx \frac{a+b}{a+b+c}\epsilon_1 + \epsilon_2\end{aligned}$$

3-5

Csci 5304 – September 16, 2013

➤ If we redid the computation as $y_2 = fl(a + fl(b+c))$ we would find

$$\begin{aligned}fl(a + fl(b+c)) &= (a+b+c)(1+\epsilon_4) \\ \epsilon_4 &\approx \frac{b+c}{a+b+c}\epsilon_1 + \epsilon_2\end{aligned}$$

➤ The first error is amplified by the factor $(a+b)/y$ in the first case and $(b+c)/y$ in the second case.

➤ In order to sum n numbers more accurately, it is better to start with the small numbers first. [However, sorting before adding is not worth the cost!]

➤ But watch out if the numbers have mixed signs!

3-6

Csci 5304 – September 16, 2013

The absolute value notation

➤ For a given vector x , $|x|$ is the vector with components $|x_i|$, i.e., $|x|$ is the component-wise absolute value of x .

➤ Similarly for matrices:

$$|A| = \{|a_{ij}|\}_{i=1,\dots,m; j=1,\dots,n}$$

➤ Obvious result. The basic inequality

$$|fl(a_{ij}) - a_{ij}| \leq \underline{u} |a_{ij}|$$

translates into

$$fl(A) = A + E \quad \text{with} \quad |E| \leq \underline{u} |A|$$

➤ $A \leq B$ means $a_{ij} \leq b_{ij}$ for all $1 \leq i \leq m; 1 \leq j \leq n$

3-7

Csci 5304 – September 16, 2013

Error Analysis: Inner product

➤ Inner products are in the innermost parts of many calculations. Their analysis is important.

Lemma: If $|\delta_i| \leq \underline{u}$ and $n\underline{u} < 1$ then

$$\prod_{i=1}^n (1 + \delta_i) = 1 + \theta_n \quad \text{where} \quad |\theta_n| \leq \frac{n\underline{u}}{1 - n\underline{u}}$$

➤ Common notation $\gamma_n \equiv \frac{n\underline{u}}{1 - n\underline{u}}$

Main result on inner products:

$$|fl(x^T y) - x^T y| \leq \gamma_n |x|^T |y|$$

➤ Absolute value notation used

3-8

Csci 5304 – September 16, 2013

- When $\gamma_n \leq 1.01n\underline{u}$ then

$$|fl(x^T y) - x^T y| \leq 1.01 n \underline{u} |x|^T |y|$$

- $\gamma_n \leq 1.01n\underline{u}$ means $[1/(1 - n\underline{u})] \leq 1.01$
- For $\underline{u} = 2.0 \times 10^{-16}$, assumption $\gamma_n \leq 1.01n\underline{u}$ holds for $n \leq 4.46 \times 10^{13}$.

- Consequence of lemma:

$$|fl(A * B) - A * B| \leq \gamma_n |A| * |B|$$

- Another way to write the result (less precise) is

$$|fl(x^T y) - x^T y| \leq n \underline{u} |x|^T |y| + O(\underline{u}^2)$$

- ✎ Prove the lemma [Hint: use induction]

- ✎ Assume you use single precision for which you have $\underline{u} = 2. \times 10^{-6}$. What is the largest n for which $\gamma_n \leq 1.01n\underline{u}$ holds? Any conclusions for the use of single precision arithmetic?

- ✎ What does the main result on inner products imply for the case when $y = x$? [Contrast the relative accuracy you get in this case vs. the general case when $y \neq x$]

Backward and forward errors

- Assume the approximation \hat{y} to $y = f(x)$ is computed with arithmetic precision ϵ . Possible analysis: find an upper bound for the **Forward** error

$$|\Delta y| = |y - \hat{y}|$$

- This is not always easy.

Alternative question: find the equivalent perturbation on the initial data (x) which produces the result \hat{y} . In other words, for what Δx do we have:

$$f(x + \Delta x) = \hat{y}$$

- The value of $|\Delta x|$ is called the backward error. An analysis to find an upper bound for $|\Delta x|$ is called **Backward error analysis**.

Example:

$$A = \begin{pmatrix} a & b \\ 0 & c \end{pmatrix} \quad B = \begin{pmatrix} d & e \\ 0 & f \end{pmatrix}$$

Consider the product: $fl(A.B) =$

$$\begin{bmatrix} (ad)(1 + \epsilon_1) & [ae(1 + \epsilon_2) + bf(1 + \epsilon_3)](1 + \epsilon_4) \\ 0 & cf(1 + \epsilon_5) \end{bmatrix}$$

with $\epsilon_i \leq \underline{u}$, for $i = 1, \dots, 5$. Result can be written as:

$$\begin{bmatrix} a & b(1 + \epsilon_3)(1 + \epsilon_4) \\ 0 & c(1 + \epsilon_5) \end{bmatrix} \begin{bmatrix} d(1 + \epsilon_1) & e(1 + \epsilon_2)(1 + \epsilon_4) \\ 0 & f \end{bmatrix}$$

- So $fl(A.B) = (A + E_A)(B + E_B)$.

- **Backward errors** E_A, E_B satisfy:

$$|E_A| \leq 2\underline{u} |A| + O(\underline{u}^2); \quad |E_B| \leq 2\underline{u} |B| + O(\underline{u}^2)$$

- When solving $Ax = b$ by Gaussian Elimination, we will see that a bound on $\|e_x\|$ such that this holds exactly

$$A(x_{\text{computed}} + e_x) = b$$

is much harder to find than bounds on $\|E_A\|$, $\|e_b\|$ such that this holds exactly

$$(A + E_A)x_{\text{computed}} = (b + e_b).$$

Note: In many instances backward errors are more meaningful than forward errors: if initial data is accurate only to 4 digits for example, then my algorithm for computing x need not have 10 digits of accuracy. A backward error of order 10^{-4} is acceptable.

- ✎ Show for any x, y , there exist $\Delta x, \Delta y$ such that

$$\begin{aligned} fl(x^T y) &= (x + \Delta x)^T y, & \text{with } |\Delta x| &\leq \gamma_n |x| \\ fl(x^T y) &= x^T (y + \Delta y), & \text{with } |\Delta y| &\leq \gamma_n |y| \end{aligned}$$

- ✎ (Continuation) Let A an $m \times n$ matrix, x an n -vector, and $y = Ax$. Show that there exist a matrix ΔA such

$$fl(y) = (A + \Delta A)x, \quad \text{with } |\Delta A| \leq \gamma_n |A|$$

- ✎ (Continuation) From the above derive a result about a column of the product of two matrices A and B . Does a similar result hold for the product AB as a whole?

Supplemental notes: Floating Point Arithmetic

In most computing systems, real numbers are represented in two parts: A mantissa and an exponent. If the representation is in the base β then:

$$x = \pm (.d_1 d_2 \cdots d_m)_\beta \beta^e$$

- $.d_1 d_2 \cdots d_m$ is a fraction in the base- β representation
- e is an integer - can be negative, positive or zero.
- Generally the form is normalized in that $d_1 \neq 0$.

Example: In base 10 (for illustration)

1. 1000.12345 can be written as

$$0.100012345_{10} \times 10^4$$

2. 0.000812345 can be written as

$$0.812345_{10} \times 10^{-3}$$

- Problem with floating point arithmetic: we have to live with limited precision.

Example: Assume that we have only 5 digits of accuracy in the mantissa and 2 digits for the exponent (excluding sign).

$$.d_1 \ d_2 \ d_3 \ d_4 \ d_5 \ | \ e_1 \ e_2$$

Let us try to add 1000.2 and 1.07

1000.2 =

.1	0	0	0	2	0	4
----	---	---	---	---	---	---

 ; 1.07 =

.1	0	7	0	0	0	1
----	---	---	---	---	---	---

First task: align decimal points. The one with smallest exponent will be (internally) rewritten so its exponent matches the largest one:

$$1.07 = 0.000107 \times 10^4$$

Second task: add mantissas:

$$\begin{array}{r} 0.10002 \\ + 0.000107 \\ \hline = 0.100127 \end{array}$$

Third task: round result. Result has 6 digits - can use only 5 so we can

➤ Chop result:

.1	0	0	1	2
----	---	---	---	---

 ;

➤ Round result:

.1	0	0	1	3
----	---	---	---	---

 ;

Fourth task: Normalize result if needed (not needed here)

result with rounding:

.1	0	0	1	3	0	4
----	---	---	---	---	---	---

 ;

 Redo the same thing with 1000.2 + 3000.8

The IEEE standard

32 bit (Single precision) :




- In binary: The leading one in mantissa does not need to be represented. One bit gained. ➤ Hidden bit.
- Largest exponent: $2^7 - 1 = 127$; Smallest: $= -126$. ['bias' of 127]

64 bit (Double precision) :



- Bias of 1023 so if c is stored exponent actual exponent is 2^{c-1023}
- $e + bias = 2047$ (all ones) = special use
- Largest exponent: 1023; Smallest = -1022.
- With hidden bit: mantissa has 53 bits represented.

 Take the number 1.0 and see what will happen if you add $1/2, 1/4, \dots, 2^{-i}$. Do not forget the hidden bit!

Hidden bit

↓ ← 52 bits →

1	1	0	0	0	0	0	0	0	0	0	0	0	e	e
1	0	1	0	0	0	0	0	0	0	0	0	0	e	e
1	0	0	1	0	0	0	0	0	0	0	0	0	e	e

.....

1	0	0	0	0	0	0	0	0	0	0	0	1	e	e
1	0	0	0	0	0	0	0	0	0	0	0	0	e	e

➤ Conclusion

$fl(1 + 2^{-52}) \neq 1$ but: $fl(1 + 2^{-53}) == 1$!!