

- Introduction – motivation
- Domain partitioning and distributed sparse matrices
- Basic algorithms: distributed Matvec
- Distributed preconditioners: additive Schwarz, multiplicative Schwarz.
- Schur complement techniques
- Graph partitioning

Introduction

- Thrust of parallel computing techniques in most applications areas.
- Programming model: Message-passing seems (MPI) dominates
- Open MP and threads for small number of processors
- Important new reality: parallel programming has penetrated the 'applications' areas [Sciences and Engineering + industry]
- Problem 1: algorithms lagging behind somewhat
- Problem 2: Message passing is painful for large applications. 'Time to solution' is high.


Domain Decomposition: A Model problem

Problem:


$$\begin{cases} \Delta u = f & \text{in } \Omega \\ u = u_\Gamma & \text{on } \Gamma = \partial\Omega. \end{cases}$$

Domain:


$$\Omega = \bigcup_{i=1}^s \Omega_i,$$



- Domain decomposition or substructuring methods attempt to solve a PDE problem (e.g.) on the entire domain from problem solutions on the subdomains Ω_i .






Discretization of domain



Coefficient Matrix

Types of mappings



(a) Vertex-based; and partitioning


(b) edge-based; and

(c) element-based

- Can adapt PDE viewpoint to general sparse matrices
- Will use the graph representation and 'vertex-based' viewpoint –


Generalization: Distributed Sparse Systems

- Simple illustration: Block assignment. Assign equation i and unknown i to a given 'process'
- Naive partitioning - won't work well in practice



- Best idea is to use the adjacency graph of A :


Vertices = $\{1, 2, \dots, n\}$;
Edges: $i \rightarrow j$ iff $a_{ij} \neq 0$



Graph partitioning problem:

- Want a partition of the vertices of the graph so that
 - (1) partitions have \sim the same sizes
 - (2) interfaces are small in size

General Partitioning of a sparse linear system



$S_1 = \{1, 2, 6, 7, 11, 12\}$:
This means equations and unknowns 1, 2, 3, 6, 7, 11, 12 are assigned to Domain 1.

$S_2 = \{3, 4, 5, 8, 9, 10, 13\}$


$S_3 = \{16, 17, 18, 21, 22, 23\}$

$S_4 = \{14, 15, 19, 20, 24, 25\}$

18-9

Csci 8314 – April 24, 2014

Alternative: Map elements / edges rather than vertices




Equations/unknowns 3, 8, 12 shared by 2 domains. From distributed sparse matrix viewpoint this is an overlap of one layer

➤ Partitioners : Metis, Chaco, Scotch, ..

➤ More recent: Zoltan, H-Metis, PaToH

18-10

Csci 8314 – April 24, 2014



18-11

Csci 8314 – April 24, 2014

➤ Standard dual objective: “minimize” communication + “balance” partition sizes

➤ Recent trend: use of hypergraphs [PaToh, Hmetis,...]


18-12

Csci 8314 – April 24, 2014

A few words about hypergraphs

- Hypergraphs are very general.. Ideas borrowed from VLSI work
- Main motivation: to better represent communication volumes when partitioning a graph. Standard models face many limitations
- Hypergraphs can better express complex graph partitioning problems and provide better solutions.
- Example: completely nonsymmetric patterns ...
- .. Even rectangular matrices

Example: $V = \{1, \dots, 9\}$ and $E = \{a, \dots, e\}$ with
 $a = \{1, 2, 3, 4\}$, $b = \{3, 5, 6, 7\}$, $c = \{4, 7, 8, 9\}$,
 $d = \{6, 7, 8\}$, and $e = \{2, 9\}$



Boolean matrix:

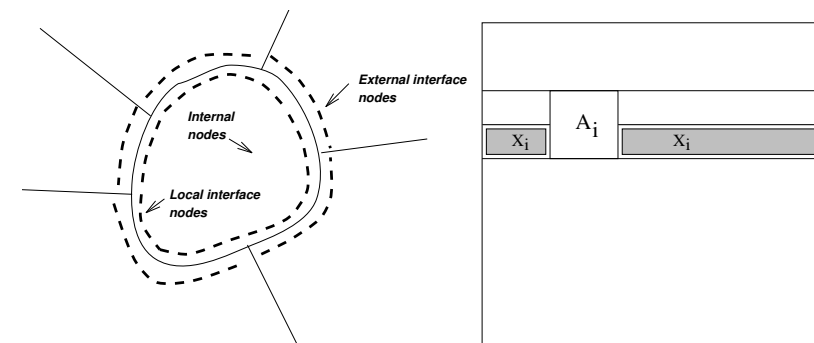
	1	2	3	4	5	6	7	8	9	
1	1	1	1	1						a
			1		1	1	1			b
				1			1	1	1	c
						1	1	1		d
	1								1	e

$A =$

Distributed Sparse matrices (continued)


- Once a good partitioning is found, questions are:
 1. How to represent this partitioning?
 2. What is a good data structure for representing distributed sparse matrices?
 3. How to set up the various "local objects" (matrices, vectors, ..)
 4. What can be done to prepare for communication that will be required during execution?

Two views of a distributed sparse matrix



- Local interface variables always ordered last.
- Need: 1) to set up the various "local objects". 2) Preprocessing to prepare for communications needed during iteration?

Local view of distributed matrix:



The local system:


$$\underbrace{\begin{pmatrix} B_i & F_i \\ E_i & C_i \end{pmatrix}}_{A_i} \begin{pmatrix} u_i \\ y_i \end{pmatrix} + \underbrace{\left(\sum_{j \in N_i} \begin{pmatrix} 0 \\ E_{ij} \end{pmatrix} y_j \right)}_{y_{ext}} = \begin{pmatrix} f_i \\ g_i \end{pmatrix}$$

- u_i : Internal variables; y_i : Interface variables

18-17

Csci 8314 – April 24, 2014

The local matrix:



The local matrix consists of 2 parts: a part (' A_{loc} ') which acts on local data and another (' B_{ext} ') which acts on remote data.

- Once the partitioning is available these parts must be identified and built locally..
- In finite elements, assembly is a local process.
- How to perform a matrix vector product? [needed by iterative schemes?]

18-18

Csci 8314 – April 24, 2014

Distributed Sparse Matrix-Vector Product Kernel

Algorithm:

1. Communicate: exchange boundary data.

Scatter x_{bound} to neighbors - Gather x_{ext} from neighbors

2. Local matrix – vector product

$$y = A_{loc} x_{loc}$$

3. External matrix – vector product

$$y = y + B_{ext} x_{ext}$$

NOTE: 1 and 2 are independent and can be overlapped.

18-19

Csci 8314 – April 24, 2014

Distributed Sparse Matrix-Vector Product

Main part of the code:

```

call MSG_bdx_send(nloc,x,y,nproc,proc,ix,ipr,ptrn,ier
c
c do local matrix-vector product for local points
c
c call amux(nloc,x,y,aloc,jaloc,ialoc)
c
c receive the boundary information
c
c call MSG_bdx_receive(nloc,x,y,nproc,proc,ix,ipr,ptrn,
c
c do local matrix-vector product for external points
c
c nrow = nloc - nbnd + 1
c call amux1(nrow,x,y(nbnd),aloc,jaloc,ialoc(nloc+1))
c
c return
    
```

18-20

Csci 8314 – April 24, 2014

The local exchange information

- List of adjacent processors (or subdomains)
- For each of these processors, lists of boundary nodes to be sent / received to / from adj. PE's.
- The receiving processor must have a matrix ordered consistently with the order in which data is received.

Requirements

- The 'set-up' routines should handle overlapping
- Should use minimal storage (only arrays of size nloc allowed).

and $e_1 = [1, 0, \dots, 0]^T$. with $\bar{H}_m = \{h_{i,j}\}_{1 \leq i \leq j+1; 1 \leq j \leq m}$.

4. **Restart:** If satisfied stop, else set $x_0 \leftarrow x_m$ and goto 1.

Distributed Flexible GMRES (FGMRES)

1. **Start:** Choose x_0 and m . Let of the Krylov subspaces. Define $\bar{H}_m \in \mathbb{R}^{(m+1) \times m}$ with $\bar{H}_m \equiv 0$. and initialize all its entries $h_{i,j}$ to zero.
2. **Arnoldi process:**
 - (a) Compute $r_0 = b - Ax_0$, $\beta = \|r_0\|_2$ and $v_1 = r_0/\beta$.
 - (b) For $j = 1, \dots, m$ do
 - Compute $z_j := M_j^{-1}v_j$; Compute $w := Az_j$;
 - For $i = 1, \dots, j$, do
 1. $h_{i,j} := (w, v_i)$
 2. $w := w - h_{i,j}v_i$
 - Compute $h_{j+1,j} = \|w\|_2$ and $v_{j+1} = w/h_{j+1,j}$.
 - (c) Define $Z_m := [z_1, \dots, z_m]$
3. **Form the approximate solution:** Compute $y_m = \operatorname{argmin}_y \|\beta e_1 - \bar{H}_m y\|_2$ and $x_m = x_0 + [z_1, z_2, \dots, z_m]y_m$

Main Operations in (F) GMRES :

1. Saxpy's – local operation – no communication
2. Dot products – global operation
3. Matrix-vector products – local operation – local communication
4. Preconditioning operations – locality varies.

Distributed Dot Product

```
/*----- call blas1 function
   tloc = DDOT(n, x, incx, y, incy);
/*----- call global reduction
   MPI_Allreduce(&tloc,&ro,1,MPI_DOUBLE,MPI_SUM,comm);
```

18-25

Csci 8314 – April 24, 2014

A remark: the global viewpoint

$$\left(\begin{array}{ccc|ccc} B_1 & & & F_1 & & \\ & B_2 & & & F_2 & \\ & & \ddots & & & \ddots \\ & & & & & & F_p \\ \hline E_1 & & & C_1 & E_{12} & \cdots & E_{1p} \\ & E_2 & & E_{21} & C_2 & \cdots & E_{2p} \\ & & \ddots & \vdots & \vdots & \ddots & \vdots \\ & & & E_p & E_{p1} & E_{p2} & \cdots & C_p \end{array} \right) \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ \vdots \\ u_p \\ y_1 \\ y_2 \\ \vdots \\ y_p \end{pmatrix} = \begin{pmatrix} f_1 \\ f_2 \\ \vdots \\ \vdots \\ f_p \\ g_1 \\ g_2 \\ \vdots \\ g_p \end{pmatrix}$$



18-26

Csci 8314 – April 24, 2014

DISTRIBUTED PRECONDITIONERS

Example: Distributed ILU(0)

➤ Global view of matrix is (for 4 processors):

$$A = \left(\begin{array}{cccc|c} A_1 & & & & F_1 \\ & A_2 & & & F_2 \\ & & A_3 & & F_3 \\ & & & A_4 & F_4 \\ \hline E_1 & E_2 & E_3 & E_4 & D \end{array} \right)$$

- A_i = local matrix restricted to internal nodes only.
- 1-st approach: Idea: ILU on this matrix – parallelism available for diagonal blocks. Define an order in which to eliminate interface unknowns.


➤ 2-nd approach: Multi-color, k -step SOR or SSOR preconditioners.

➤ 3-rd approach: Solve equations for all interface points [Schur Complement approach] – to precondition, use ideas from DD.

Example: Distributed ILU(0) – cont.


- Easy to understand from a local view of distributed matrix
- Start by selecting an order [or a “schedule”, or a “priority rule”] in which to process globally
- Then locally:
 1. Eliminate internal rows
 2. Receive rows needed to process local interface rows
 3. Process local interface rows
 4. Send local interface rows to processors needing them

A distributed view of ILU(0) – schedule based on PE numbers



Note: any schedule can be used provided neighbors have different labels. Example: can use coloring.

A distributed view of ILU(0) – schedule based on PE coloring



- See [S. Ma and YS. 1994]
- Generalized ILU(k): D. Hysom and A. Pothen '00.
- Used in pARMS for preconditioning Schur complement


Domain Decomposition–Type preconditioners

- Schwarz Preconditioners
- Schur-complement based Preconditioners
- Multi-level ILU-type Preconditioners

► Observation: Often, in practical applications, Schwarz Preconditioners are used : SUB-OPTIMAL

Domain-Decomposition Preconditioners (cont.)

Local view of distributed matrix:



Block Jacobi Iteration (Additive Schwarz):

1. Obtain external data y_i
2. Compute (update) local residual

$$r_i = (b - Ax)_i = b_i - A_i x_i - B_i y_i$$
3. Solve $A_i \delta_i = r_i$
4. Update solution $x_i = x_i + \delta_i$

► Multiplicative Schwarz. Need a coloring of the subdomains so that:

► No two adjacent subdomains share same color


Multicolor Block SOR Iteration (Multiplicative Schwarz):

1. Do $col = 1, \dots, numcols$
2. If ($col.eq.mycol$) Then
3. Obtain external data y_i
4. Update local residual $r_i = (b - Ax)_i$
5. Solve $A_i \delta_i = r_i$
6. Update solution $x_i = x_i + \delta_i$
7. EndIf
8. EndDo


Breaking the sequential color loop

► “Color” loop is sequential. Can be broken in several different ways.

(1) Have a few subdomains per processors



(2) Separate interior nodes from interface nodes (2-level blocking)



(3) Use a block-GMRES algorithm - with Block-size = number of colors. SOR step targets a different color on each column of the block ➤ no idle time.

Local Solves

➤ Each local system $A_i \delta_i = r_i$ can be solved in three ways:

1. By a (sparse) direct solver
2. Using a standard preconditioned Krylov solver
3. Doing a backward-forward solution associated with an accurate ILU (e.g. ILUT) preconditioner


➤ We only use (2) with a small number of inner steps (up to 10) or (3).

SCHUR COMPLEMENT-BASED PRECONDITIONERS

Schur complement system

Local system can be written as

$$A_i x_i + X_i y_{i,ext} = b_i. \quad (1)$$



x_i = vector of local unknowns, $y_{i,ext}$ = external interface variables, and b_i = local part of RHS.

➤ Local equations

$$\begin{pmatrix} B_i & F_i \\ E_i & C_i \end{pmatrix} \begin{pmatrix} u_i \\ y_i \end{pmatrix} + \begin{pmatrix} 0 \\ \sum_{j \in N_i} E_{ij} y_j \end{pmatrix} = \begin{pmatrix} f_i \\ g_i \end{pmatrix} \quad (2)$$

➤ eliminate u_i from the above system:

$$S_i y_i + \sum_{j \in N_i} E_{ij} y_j = g_i - E_i B_i^{-1} f_i \equiv g'_i,$$

where S_i is the “local” Schur complement

$$S_i = C_i - E_i B_i^{-1} F_i. \quad (3)$$

Structure of Schur complement system

Global Schur complement system: $Sy = g'$ with :

$$S = \begin{pmatrix} S_1 & E_{12} & \dots & E_{1p} \\ E_{21} & S_2 & \dots & E_{2p} \\ \vdots & & \ddots & \vdots \\ E_{p1} & E_{p-1,2} & \dots & S_p \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_p \end{pmatrix} = \begin{pmatrix} g'_1 \\ g'_2 \\ \vdots \\ g'_p \end{pmatrix}.$$

- E_{ij} 's are sparse = same as in the original matrix
- Can solve global Schur complement system iteratively. Back-substitute to recover rest of variables (internal).
- Can use the procedure as a preconditioning to global system.

Simplest idea: Schur Complement Iterations

$\begin{pmatrix} u_i \\ y_i \end{pmatrix}$ Internal variables
Interface variables

- Do a global primary iteration (e.g., block-Jacobi)
- Then accelerate only the y variables (with a Krylov method)

Still need to precondition..

Approximate Schur-LU

- Two-level method based on induced preconditioner. Global system can also be viewed as

$$\begin{pmatrix} B & F \\ E & C \end{pmatrix} \begin{pmatrix} u \\ y \end{pmatrix} = \begin{pmatrix} f \\ g \end{pmatrix}, \quad B = \left(\begin{array}{ccc|c} B_1 & & & F_1 \\ & B_2 & & F_2 \\ & & \ddots & \vdots \\ & & & B_p & F_p \\ \hline E_1 & E_2 & \dots & E_p & C \end{array} \right)$$

Block LU factorization of A :

$$\begin{pmatrix} B & F \\ E & C \end{pmatrix} = \begin{pmatrix} B & 0 \\ E & S \end{pmatrix} \begin{pmatrix} I & B^{-1}F \\ 0 & I \end{pmatrix},$$

Preconditioning:

$$L = \begin{pmatrix} B & 0 \\ E & M_S \end{pmatrix} \quad \text{and} \quad U = \begin{pmatrix} I & B^{-1}F \\ 0 & I \end{pmatrix}$$

with M_S = some approximation to S .

- Preconditioning to global system can be induced from any preconditioning on Schur complement.

Rewrite local Schur system as

$$y_i + S_i^{-1} \sum_{j \in N_i} E_{ij} y_j = S_i^{-1} [g_i - E_i B_i^{-1} f_i].$$

- equivalent to Block-Jacobi preconditioner for Schur complement.
- Solve with, e.g., a few s (e.g., 5) of GMRES

GRAPH PARTITIONING

- Question: How to solve with S_i ?

- Can use LU factorization of local matrix $A_i = \begin{pmatrix} B_i & F_i \\ E_i & C_i \end{pmatrix}$

and exploit the relation:

$$A_i = \begin{pmatrix} L_{B_i} & 0 \\ E_i U_{B_i}^{-1} & L_{S_i} \end{pmatrix} \begin{pmatrix} U_{B_i} & L_{B_i}^{-1} F_i \\ 0 & U_{S_i} \end{pmatrix} \rightarrow L_{S_i} U_{S_i} = S_i$$

- Need only the (l) LU factorization of the A_i [rest is already available]
- Very easy implementation of (parallel) Schur complement techniques for vertex-based partitioned systems : YS-Sosonkina '97; YS-Sosonkina-Zhang '99.

Three approaches to graph partitioning:

1. Spectral methods (Recursive Spectral Bisection)
2. Geometric techniques [Houstis & Rice et al., Miller, Vavasis, Teng et al.] Coordinates are required.
3. Graph Theory techniques [use graph, but no coordinates]
 - Currently best known technique is Metis (multi-level algorithm)
 - Simplest idea: Recursive Graph Bisection; Nested dissection (George & Liu, 1980; Liu 1992)
 - Advantages: simplicity – no coordinates required

Background: Graph Laplaceans

- “Laplace-type” matrices associated with general undirected graphs – useful in many applications
- Given a graph $G = (V, E)$ define
 - A matrix W of weights w_{ij} for each edge
 - Assume $w_{ij} \geq 0$, $w_{ii} = 0$, and $w_{ij} = w_{ji} \forall (i, j)$
 - The diagonal matrix $D = \text{diag}(d_i)$ with $d_i = \sum_{j \neq i} w_{ij}$
- Corresponding **graph Laplacean** of G is:

$$L = D - W$$

- Gershgorin's theorem $\rightarrow L$ is positive semidefinite

18-49

Csci 8314 – April 24, 2014

Graph Laplaceans: Another viewpoint

- Let N be the incidence matrix: $N_{ij} = \pm 1$ if i -th edge is incident on the j -th vertex.
- For example: $A \leftrightarrow C, D$, $B \leftrightarrow D$, $C \leftrightarrow A$, $D \leftrightarrow A, B$ (undirected graph):

$$N = \begin{bmatrix} 1 & 0 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & -1 & 0 & 1 \end{bmatrix},$$

yielding Laplacian = diagonal matrix of degrees – Adjacency matrix

$$N^T N = L = \begin{bmatrix} 2 & 0 & -1 & -1 \\ 0 & 1 & 0 & -1 \\ -1 & 0 & 1 & 0 \\ -1 & -1 & 0 & 2 \end{bmatrix}.$$

18-50

Csci 8314 – April 24, 2014

Spectral Graph Partitioning

- Consider simplest case:

$$w_{ij} = \begin{cases} 1 & \text{if } (i, j) \in E \text{ and } i \neq j \\ 0 & \text{else} \end{cases} \quad D = \text{diag} \left[d_i = \sum_{j \neq i} w_{ij} \right]$$

Property: (Graph partitioning) Consider situation when $w_{ij} \in \{0, 1\}$. If x is a vector of signs (± 1) then

$$x^T L x = 4 \times (\text{'number of edge cuts'})$$

edge-cut = pair (i, j) with $x_i \neq x_j$

18-51

Csci 8314 – April 24, 2014

- Consequence: Can be used for partitioning graphs, or ‘clustering’ [take $p = \text{sign}(u_2)$, where u_2 = 2nd smallest eigenvector..]
- Would like to minimize (Lx, x) subject to $x \in \{-1, 1\}^n$ and $e^T x = 0$ [balanced sets]
- Will solve a relaxed form of this problem

Background: Consider any symmetric (real) matrix A with eigenvalues $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$ and eigenvectors u_1, \dots, u_n

- Recall that:
(Min reached for $x = u_1$)

$$\min_{x \in \mathbb{R}^n} \frac{(Ax, x)}{(x, x)} = \lambda_1$$

- In addition:
(Min reached for $x = u_2$)

$$\min_{x \perp u_1} \frac{(Ax, x)}{(x, x)} = \lambda_2$$

- For a graph Laplacean $u_1 = e$ = vector of all ones and
- ...vector u_2 is called the Fiedler vector. It solves a **relaxed** form of the problem -

$$\min_{x \in \{-1,1\}^n; e^T x = 0} \frac{(Lx, x)}{(x, x)} \rightarrow \min_{x \in \mathbb{R}^n; e^T x = 0} \frac{(Lx, x)}{(x, x)}$$

- Define $v = u_2$ then $lab = \text{sign}(v - \text{med}(v))$

Normalized Graph Cuts

Mark a partitioning of the vertices: $n_- = 1, n_+ = 3$
 $v = [1, 1, 1, -3]^T / \sqrt{3 \cdot 1} = [n_-, n_-, n_-, -n_+]^T / \sqrt{n_- n_+}$.

Then

$$\frac{v^T L v}{v^T v} = |\text{cut}| \cdot \left(\frac{1}{n_-} + \frac{1}{n_+} \right)$$

and $v^T e = 0$, where $e = [1, 1, 1, 1]^T$ = eigenvector of L .

- Approximately minimize this with an eigenvector of L :
 -1.E-15 (.500000 .500000 .500000 .500000) ← 'null' vector
 .585786 (-.27059 .653281 -.65328 .270598) ← 'Fiedler' vector
 2.00000 (.500000 -.50000 -.50000 .500000)
 3.41421 (.653281 .270598 -.27059 -.65328)

The Level Set Expansion Algorithm

- Given: p nodes 'uniformly' spread in the graph (roughly same distance from one another).

Do a level-set traversal from each node simultaneously.

Best described for an example on a 15×15 five – point Finite Difference grid.

- See [Goehring-Saad '94, See Cai-Saad '95]

